

LAB 13

ABDUL RAUF

2023-02-05

Contents

1	RIDGE REGRESSION:	2
2	Load the Data	2
3	Fit the ridge regression model	2
4	Performing k-fold cross-validation	3
5	Plot the test MSE's vs. lambda values	3
6	Coefficients of best model	3
7	Prediction for the response value of a new observation	4
8	Ridge trace plot	4
9	Find R-squared of model on training data	5

```
__NAME__:-**ABDUL RAUF** \
__ENRL NO__:-**GL6092** \
__ROLL NO__:-**22DSMSA116** \
```

1 RIDGE REGRESSION:

Ridge regression is a model tuning method that is used to analyse any data that suffers from multicollinearity. This method performs L2 regularization. When the issue of multicollinearity occurs, least-squares are unbiased, and variances are large, this results in predicted values being far away from the actual values.

To perform lasso regression, we'll use functions from the `glmnet()` package. This package requires the response variable to be a vector and the set of predictor variables to be of the class `data.matrix`.

use the `glmnet()` function to fit the ridge regression model and specify `alpha=0`.

Note that setting `alpha` equal to 0 is equivalent to using ridge regression and setting `alpha` to some value between 0 and 1 is equivalent to using an elastic net.

To determine what value to use for `lambda`, we'll perform k-fold cross-validation and identify the `lambda` value that produces the lowest test mean squared error (MSE). `cv.glmnet()` automatically performs k-fold cross validation using `k = 10` folds.

2 Load the Data

```
library(glmnet)
library(ISLR2)

#Define predictor and response variables
y <- mtcars$hp
x <- data.matrix(mtcars[, c('mpg', 'wt', 'drat', 'qsec')])
```

3 Fit the ridge regression model

```
model <- glmnet(x, y, alpha = 0)
summary(model)
```

```
##           Length Class      Mode
## a0         100    -none-   numeric
## beta        400  dgCMatrix S4
## df          100    -none-   numeric
## dim           2    -none-   numeric
## lambda       100    -none-   numeric
## dev.ratio  100    -none-   numeric
## nulldev       1    -none-   numeric
## npasses       1    -none-   numeric
## jerr          1    -none-   numeric
## offset        1    -none-   logical
## call          4    -none-    call
## nobs         1    -none-   numeric
```

4 Performing k-fold cross-validation

```
cv_model <- cv.glmnet(x, y, alpha = 0)
best_lambda <- cv_model$lambda.min


#display optimal lambda value
best_lambda
```

```
## [1] 7.599184
```

The lambda value that minimizes the test MSE turns out to be 11.02511.

5 Plot the test MSE's vs. lambda values

```
plot(cv_model)
```



6 Coefficients of best model

```
best_model <- glmnet(x, y, alpha = 0, lambda = best_lambda)
coef(best_model)
```

```
## 5 x 1 sparse Matrix of class "dgCMatrix"
##                s0
## (Intercept) 477.88658049
## mpg        -3.29585113
## wt         20.32422537
## drat       -0.09481675
## qsec       -18.49039687
```

the coefficient of the `drat` becomes exactly equal to 0 ,it get completely dropped the model.

7 Prediction for the response value of a new observation


```
new = matrix(c(24, 2.8, 3.4, 18.4), nrow=1, ncol=4)
predict(best_model, s = best_lambda, newx = new)
```

```
##                s1
## [1,] 115.1483
```

based on the new input value as a test data , model predict that car have an hp value of 115.2554.

8 Ridge trace plot

```
plot(model,xvar = "lambda")
```



all the variables are approaching towards 0.

9 Find R-squared of model on training data

```
y_predicted <- predict(best_model, s = best_lambda, newx = x)

sst <- sum((y - mean(y))^2)
sse <- sum((y_predicted - y)^2)

rsq <- 1 - sse/sst
rsq
```

```
## [1] 0.8025993
```

R-squared turns out to be 0.7987891 which means 79.87% of the variation in the response is explained by model.